

✓ Detecção de Casos de COVID-19 em Tomografias (versão PyTorch)

Agenda

- Carregar dataset (ImageFolder)
- Pré-processamento e augmentations
- Transfer Learning (ResNet18)
- Treino, validação, checkpoint, early stopping
- Avaliação final e plots

```
!pip install kaggle
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.12/dist-packages (1.7.4.5)
Requirement already satisfied: bleach in /usr/local/lib/python3.12/dist-packages (from kaggle) (6.3.0)
Requirement already satisfied: certifi>=14.05.14 in /usr/local/lib/python3.12/dist-packages (from kaggle) (2025.11.12)
Requirement already satisfied: charset-normalizer in /usr/local/lib/python3.12/dist-packages (from kaggle) (3.4.4)
Requirement already satisfied: idna in /usr/local/lib/python3.12/dist-packages (from kaggle) (3.11)
Requirement already satisfied: protobuf in /usr/local/lib/python3.12/dist-packages (from kaggle) (5.29.5)
Requirement already satisfied: python-dateutil>=2.5.3 in /usr/local/lib/python3.12/dist-packages (from kaggle) (2.9.0.post0)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.12/dist-packages (from kaggle) (8.0.4)
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages (from kaggle) (2.32.4)
Requirement already satisfied: setuptools>=21.0.0 in /usr/local/lib/python3.12/dist-packages (from kaggle) (75.2.0)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.12/dist-packages (from kaggle) (1.17.0)
Requirement already satisfied: text-unidecode in /usr/local/lib/python3.12/dist-packages (from kaggle) (1.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from kaggle) (4.67.1)
Requirement already satisfied: urllib3>=1.15.1 in /usr/local/lib/python3.12/dist-packages (from kaggle) (2.5.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.12/dist-packages (from kaggle) (0.5.1)
```

```
from google.colab import files
files.upload() # Envie o arquivo kaggle.json
```

Escolher arquivos Nenhum arquivo escolhido Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving archive (3).zip to archive (3).zip

```
# Cell: Imports e configuração
import os
import random

from pathlib import Path
import numpy as np
import matplotlib.pyplot as plt
from tqdm import tqdm

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, random_split
from torchvision import transforms, datasets, models

from sklearn.metrics import classification_report, confusion_matrix

# Configurações
SEED = 42
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed_all(SEED)

DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", DEVICE)
```

Device: cuda

```
# Paths e hiper-parâmetros (ajuste aqui)
# DATA_DIR = "/kaggle/input" # <- ajuste para o caminho onde você extraiu o dataset
# The dataset was unzipped to /content/covid_dataset
DATA_DIR = "/content/covid_dataset"
# Espera-se uma estrutura: DATA_DIR/class1/*.png, DATA_DIR/class2/*.png (ImageFolder)
BATCH_SIZE = 16
IMG_SIZE = 224
NUM_EPOCHS = 25
LR = 2e-4
NUM_WORKERS = 4
```

```
PATIENCE = 6 # early stopping
MODEL_SAVE_PATH = "best_model_resnet18.pt"
```

```
import zipfile
import os

# Path to the uploaded zip file
zip_path = '/content/archive (3).zip'
# Directory where the dataset will be extracted
extract_base_path = '/content/covid_dataset'

# Create the extraction directory if it doesn't exist
if not os.path.exists(extract_base_path):
    os.makedirs(extract_base_path)

# Unzip the file
print(f"Unzipping {zip_path} to {extract_base_path}...")
with zipfile.ZipFile(zip_path, 'r') as zip_ref:
    zip_ref.extractall(extract_base_path)

print("Unzipping complete.")

# Optional: List contents to verify structure and find the actual dataset root
print("Contents of extracted directory:")
for root, dirs, files in os.walk(extract_base_path):
    level = root.replace(extract_base_path, '').count(os.sep)
    indent = ' ' * 4 * (level)
    print(f'{indent}{os.path.basename(root)}/')
    subindent = ' ' * 4 * (level + 1)
    for f in files:
        if not f.startswith('.'): # Ignore hidden files
            print(f'{subindent}{f}')
    if level > 1: # Limit depth for listing to avoid long output
        break

# The actual DATA_DIR should be adjusted based on the unzipped structure.
# Assuming the zip extracts into a folder named 'CT_Scan_Dataset' like:
# /content/covid_dataset/CT_Scan_Dataset/COVID
# /content/covid_dataset/CT_Scan_Dataset/Non-COVID
# The DATA_DIR variable is updated in cell 'nPLXoMoN7YCI' to reflect this.
```

Unzipping /content/archive (3).zip to /content/covid_dataset...

Unzipping complete.

Contents of extracted directory:

```
covid_dataset/
non-COVID/
    Non-Covid (1133).png
    Non-Covid (399).png
    Non-Covid (145).png
    Non-Covid (410).png
    Non-Covid (624).png
    Non-Covid (1054).png
    Non-Covid (1210).png
    Non-Covid (202).png
    Non-Covid (621).png
    Non-Covid (458).png
    Non-Covid (237).png
    Non-Covid (850).png
    Non-Covid (17).png
    Non-Covid (669).png
    Non-Covid (14).png
    Non-Covid (169).png
    Non-Covid (1199).png
    Non-Covid (710).png
    Non-Covid (421).png
    Non-Covid (497).png
    Non-Covid (1177).png
    Non-Covid (451).png
    Non-Covid (10).png
    Non-Covid (531).png
    Non-Covid (46).png
    Non-Covid (157).png
    Non-Covid (299).png
    Non-Covid (348).png
    Non-Covid (1132).png
    Non-Covid (527).png
    Non-Covid (358).png
    Non-Covid (1219).png
    Non-Covid (622).png
    Non-Covid (709).png
    Non-Covid (320).png
    Non-Covid (909).png
    Non-Covid (1172).png
    Non-Covid (774).png
    Non-Covid (191).png
```

```
Non-Covid (1156).png
Non-Covid (1023).png
Non-Covid (336).png
Non-Covid (288).png
Non-Covid (1134).png
Non-Covid (113).png
Non-Covid (532).png
Non-Covid (271).png
Non-Covid (121).png
Non-Covid (420).png
Non-Covid (227).png
Non-Covid (18).png
Non-Covid (181).png
```

```
# Augmentations e transforms
train_transform = transforms.Compose([
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(10),
    transforms.ColorJitter(brightness=0.1, contrast=0.1),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485,0.456,0.406], std=[0.229,0.224,0.225])
])

val_transform = transforms.Compose([
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485,0.456,0.406], std=[0.229,0.224,0.225])
])

# Carregar todo dataset via ImageFolder
full_dataset = datasets.ImageFolder(root=DATA_DIR, transform=train_transform)
class_names = full_dataset.classes
num_classes = len(class_names)
print("Classes:", class_names, "Num classes:", num_classes)

# Split train / val (ex.: 80/20)
val_ratio = 0.2
n_total = len(full_dataset)
n_val = int(n_total * val_ratio)
n_train = n_total - n_val

train_dataset, val_dataset = random_split(
    full_dataset,
    [n_train, n_val],
    generator=torch.Generator().manual_seed(SEED)
)

# Aplica a transformação de validação ao conjunto de validação
# Isso é necessário porque `random_split` herda a transformação do `full_dataset`
# que é `train_transform`. Precisamos aplicar `val_transform` ao `val_dataset`.
val_dataset.dataset.transform = val_transform

print(f"Tamanho train: {len(train_dataset)}, val: {len(val_dataset)}")

train_loader = DataLoader(
    train_dataset,
    batch_size=BATCH_SIZE,
    shuffle=True,
    num_workers=NUM_WORKERS,
    pin_memory=True
)

val_loader = DataLoader(
    val_dataset,
    batch_size=BATCH_SIZE,
    shuffle=False,
    num_workers=NUM_WORKERS,
    pin_memory=True
)

Classes: ['COVID', 'non-COVID'] Num classes: 2
Tamanho train: 1985, val: 496
/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:627: UserWarning: This DataLoader will create 4 workers
warnings.warn()
```

```
# Model (Transfer Learning com ResNet18)
def build_model(num_classes):
    model = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)
    # Congelar todos os parâmetros do modelo pré-treinado
    for param in model.parameters():
        param.requires_grad = False
```

```

# Substituir a camada final (fully connected) para nosso número de classes
num_ftrs = model.fc.in_features
model.fc = nn.Linear(num_ftrs, num_classes)
model = model.to(DEVICE)
return model

model = build_model(num_classes)
print(model)

# Loss e Optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.fc.parameters(), lr=LR) # Apenas camada final treinável

# Scheduler para ajuste do Learning Rate
scheduler = optim.lr_scheduler.ReduceLROnPlateau(
    optimizer,
    mode='min',
    factor=0.1,
    patience=PATIENCE // 2
)

Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to /root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth
100%|██████████| 44.7M/44.7M [00:00<00:00, 228MB/s]
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer3): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
  )
)
```

```

# Funções de treino e validação
def train_epoch(model, dataloader, criterion, optimizer, device):
    model.train()
    running_loss = 0.0
    correct_predictions = 0
    total_predictions = 0

```

```

for inputs, labels in tqdm(dataloader, desc="Training"):
    inputs, labels = inputs.to(device), labels.to(device)

    optimizer.zero_grad()
    outputs = model(inputs)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    running_loss += loss.item() * inputs.size(0)
    _, predicted = torch.max(outputs.data, 1)
    total_predictions += labels.size(0)
    correct_predictions += (predicted == labels).sum().item()

epoch_loss = running_loss / len(dataloader.dataset)
epoch_acc = correct_predictions / total_predictions
return epoch_loss, epoch_acc

def validate_epoch(model, dataloader, criterion, device):
    model.eval()
    running_loss = 0.0
    correct_predictions = 0
    total_predictions = 0
    with torch.no_grad():
        for inputs, labels in tqdm(dataloader, desc="Validation"):
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            loss = criterion(outputs, labels)

            running_loss += loss.item() * inputs.size(0)
            _, predicted = torch.max(outputs.data, 1)
            total_predictions += labels.size(0)
            correct_predictions += (predicted == labels).sum().item()

    epoch_loss = running_loss / len(dataloader.dataset)
    epoch_acc = correct_predictions / total_predictions
    return epoch_loss, epoch_acc

# Loop de Treino
best_val_loss = float('inf')
epochs_no_improve = 0

history = {'train_loss': [], 'train_acc': [], 'val_loss': [], 'val_acc': []}

print("Starting training...")
for epoch in range(NUM_EPOCHS):
    print(f"Epoch {epoch+1}/{NUM_EPOCHS}")

    train_loss, train_acc = train_epoch(model, train_loader, criterion, optimizer, DEVICE)
    val_loss, val_acc = validate_epoch(model, val_loader, criterion, DEVICE)

    scheduler.step(val_loss) # Atualiza o learning rate com base na val_loss

    history['train_loss'].append(train_loss)
    history['train_acc'].append(train_acc)
    history['val_loss'].append(val_loss)
    history['val_acc'].append(val_acc)

    print(f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f}")
    print(f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f}")

    # Checkpoint e Early Stopping
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        epochs_no_improve = 0
        torch.save(model.state_dict(), MODEL_SAVE_PATH)
        print("Model saved! Validation loss improved.")
    else:
        epochs_no_improve += 1
        print(f"Validation loss did not improve for {epochs_no_improve} epoch(s).")
        if epochs_no_improve == PATIENCE:
            print(f"Early stopping triggered after {epoch+1} epochs.")
            break

print("Training complete.")

```

Starting training...
Epoch 1/25
Training: 100%|██████████| 125/125 [00:11<00:00, 11.01it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 10.74it/s]
Train Loss: 0.6174, Train Acc: 0.6650

```

Val Loss: 0.5535, Val Acc: 0.7258
Model saved! Validation loss improved.
Epoch 2/25
Training: 100%|██████████| 125/125 [00:10<00:00, 11.62it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 11.92it/s]
Train Loss: 0.5246, Train Acc: 0.7476
Val Loss: 0.4785, Val Acc: 0.7843
Model saved! Validation loss improved.
Epoch 3/25
Training: 100%|██████████| 125/125 [00:10<00:00, 11.83it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 12.14it/s]
Train Loss: 0.4648, Train Acc: 0.8091
Val Loss: 0.4396, Val Acc: 0.8125
Model saved! Validation loss improved.
Epoch 4/25
Training: 100%|██████████| 125/125 [00:10<00:00, 11.74it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 11.85it/s]
Train Loss: 0.4297, Train Acc: 0.8081
Val Loss: 0.4083, Val Acc: 0.8306
Model saved! Validation loss improved.
Epoch 5/25
Training: 100%|██████████| 125/125 [00:10<00:00, 11.93it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 11.97it/s]
Train Loss: 0.4118, Train Acc: 0.8302
Val Loss: 0.3915, Val Acc: 0.8387
Model saved! Validation loss improved.
Epoch 6/25
Training: 100%|██████████| 125/125 [00:10<00:00, 11.71it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 12.01it/s]
Train Loss: 0.3956, Train Acc: 0.8252
Val Loss: 0.3704, Val Acc: 0.8488
Model saved! Validation loss improved.
Epoch 7/25
Training: 100%|██████████| 125/125 [00:10<00:00, 11.87it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 12.07it/s]
Train Loss: 0.3698, Train Acc: 0.8453
Val Loss: 0.3637, Val Acc: 0.8589
Model saved! Validation loss improved.
Epoch 8/25
Training: 100%|██████████| 125/125 [00:10<00:00, 11.92it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 12.07it/s]
Train Loss: 0.3671, Train Acc: 0.8458
Val Loss: 0.3478, Val Acc: 0.8508
Model saved! Validation loss improved.
Epoch 9/25
Training: 100%|██████████| 125/125 [00:10<00:00, 11.73it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 11.72it/s]
Train Loss: 0.3618, Train Acc: 0.8524
Val Loss: 0.3726, Val Acc: 0.8327
Validation loss did not improve for 1 epoch(s).
Epoch 10/25
Training: 100%|██████████| 125/125 [00:10<00:00, 11.94it/s]
Validation: 100%|██████████| 31/31 [00:02<00:00, 12.15it/s]

```

```

model = models.resnet18(pretrained=True)
# Substituir a última camada fully-connected
in_features = model.fc.in_features
model.fc = nn.Linear(in_features, num_classes)

model = model.to(DEVICE)

# Se quiser "congelar" as camadas iniciais, descomente:
# for param in model.parameters():
#     param.requires_grad = False
# for param in model.fc.parameters():
#     param.requires_grad = True

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=LR)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.5, patience=2)

/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated
  warnings.warn(
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or
  warnings.warn(msg)

```

```

def train_one_epoch(model, loader, criterion, optimizer, device):
    model.train()
    running_loss = 0.0
    running_corrects = 0
    total = 0

    for inputs, labels in tqdm(loader, desc="Train", leave=False):
        inputs = inputs.to(device)
        labels = labels.to(device)

```

```

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * inputs.size(0)
        _, preds = torch.max(outputs, 1)
        running_corrects += torch.sum(preds == labels.data).item()
        total += inputs.size(0)

    epoch_loss = running_loss / total
    epoch_acc = running_corrects / total
    return epoch_loss, epoch_acc

def validate(model, loader, criterion, device):
    model.eval()
    running_loss = 0.0
    running_corrects = 0
    total = 0
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for inputs, labels in tqdm(loader, desc="Val", leave=False):
            inputs = inputs.to(device)
            labels = labels.to(device)
            outputs = model(inputs)
            loss = criterion(outputs, labels)

            running_loss += loss.item() * inputs.size(0)
            _, preds = torch.max(outputs, 1)
            running_corrects += torch.sum(preds == labels.data).item()
            total += inputs.size(0)

            all_preds.extend(preds.cpu().numpy().tolist())
            all_labels.extend(labels.cpu().numpy().tolist())

    epoch_loss = running_loss / total
    epoch_acc = running_corrects / total
    return epoch_loss, epoch_acc, all_preds, all_labels

```

```

best_val_loss = float('inf')
best_val_acc = 0.0
patience_counter = 0

history = {
    'train_loss': [], 'train_acc': [],
    'val_loss': [], 'val_acc': []
}

for epoch in range(1, NUM_EPOCHS + 1):
    print(f"Epoch {epoch}/{NUM_EPOCHS}")
    train_loss, train_acc = train_one_epoch(model, train_loader, criterion, optimizer, DEVICE)
    val_loss, val_acc, _, _ = validate(model, val_loader, criterion, DEVICE)
    scheduler.step(val_loss) # opcional

    history['train_loss'].append(train_loss)
    history['train_acc'].append(train_acc)
    history['val_loss'].append(val_loss)
    history['val_acc'].append(val_acc)

    print(f" train_loss: {train_loss:.4f} train_acc: {train_acc:.4f}")
    print(f" val_loss: {val_loss:.4f} val_acc: {val_acc:.4f}")

    # Checkpoint by val_loss (mudar se preferir val_acc)
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        best_val_acc = val_acc
        torch.save({
            'model_state_dict': model.state_dict(),
            'optimizer_state_dict': optimizer.state_dict(),
            'class_names': class_names
        }, MODEL_SAVE_PATH)
        print(" Saved best model.")
        patience_counter = 0
    else:
        patience_counter += 1
        print(f" Patience: {patience_counter}/{PATIENCE}")

    if patience_counter >= PATIENCE:

```

```

        print("Early stopping triggered.")
        break

    print("Treino finalizado. Best val_loss:", best_val_loss, "best val_acc:", best_val_acc)

Epoch 1/25
Train:  0% | 0/125 [00:00<?, ?it/s]/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:627:
         warnings.warn(
train_loss: 0.2492 train_acc: 0.8882
val_loss:  0.1705 val_acc:   0.9315
Saved best model.
Epoch 2/25
train_loss: 0.0486 train_acc: 0.9864
val_loss:  0.0680 val_acc:   0.9778
Saved best model.
Epoch 3/25
train_loss: 0.0731 train_acc: 0.9768
val_loss:  0.0645 val_acc:   0.9758
Saved best model.
Epoch 4/25
train_loss: 0.0614 train_acc: 0.9814
val_loss:  0.0540 val_acc:   0.9698
Saved best model.
Epoch 5/25
train_loss: 0.0471 train_acc: 0.9839
val_loss:  0.0737 val_acc:   0.9758
Patience: 1/6
Epoch 6/25
train_loss: 0.0325 train_acc: 0.9899
val_loss:  0.0836 val_acc:   0.9738
Patience: 2/6
Epoch 7/25
train_loss: 0.0166 train_acc: 0.9960
val_loss:  0.0249 val_acc:   0.9919
Saved best model.
Epoch 8/25
train_loss: 0.0053 train_acc: 0.9980
val_loss:  0.0266 val_acc:   0.9899
Patience: 1/6
Epoch 9/25
train_loss: 0.0033 train_acc: 0.9995
val_loss:  0.0095 val_acc:   0.9980
Saved best model.
Epoch 10/25
train_loss: 0.0015 train_acc: 1.0000
val_loss:  0.0180 val_acc:   0.9919
Patience: 1/6
Epoch 11/25
train_loss: 0.0313 train_acc: 0.9909
val_loss:  0.0472 val_acc:   0.9819
Patience: 2/6
Epoch 12/25
train_loss: 0.0418 train_acc: 0.9869
val_loss:  0.2387 val_acc:   0.9214
Patience: 3/6
Epoch 13/25
train_loss: 0.0090 train_acc: 0.9975
val_loss:  0.0301 val_acc:   0.9919
Patience: 4/6
Epoch 14/25
train_loss: 0.0062 train_acc: 0.9975
val_loss:  0.0215 val_acc:   0.9919
^ ^ ^

```

```

# Carregar melhor checkpoint
checkpoint = torch.load(MODEL_SAVE_PATH, map_location=DEVICE)
model.load_state_dict(checkpoint['model_state_dict'])
model.to(DEVICE)
model.eval()

# Rodar validação final para obter previsões
val_loss, val_acc, preds, labels = validate(model, val_loader, criterion, DEVICE)
print("Val loss:", val_loss, "Val acc:", val_acc)

# Relatório
print("\nClassification Report:")
print(classification_report(labels, preds, target_names=class_names))

# Matriz de confusão
cm = confusion_matrix(labels, preds)
print("Confusion Matrix:\n", cm)

```

Val loss: 0.009547496510241482 Val acc: 0.9979838709677419

Classification Report:

	precision	recall	f1-score	support
COVID	1.00	1.00	1.00	231
non-COVID	1.00	1.00	1.00	265
accuracy			1.00	496
macro avg	1.00	1.00	1.00	496
weighted avg	1.00	1.00	1.00	496

Confusion Matrix:

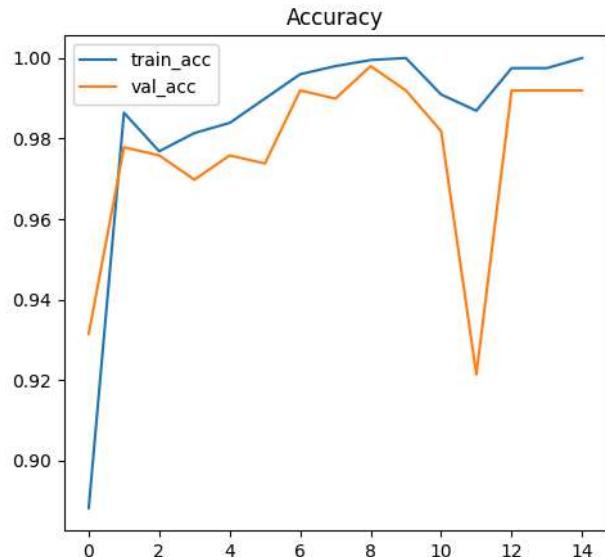
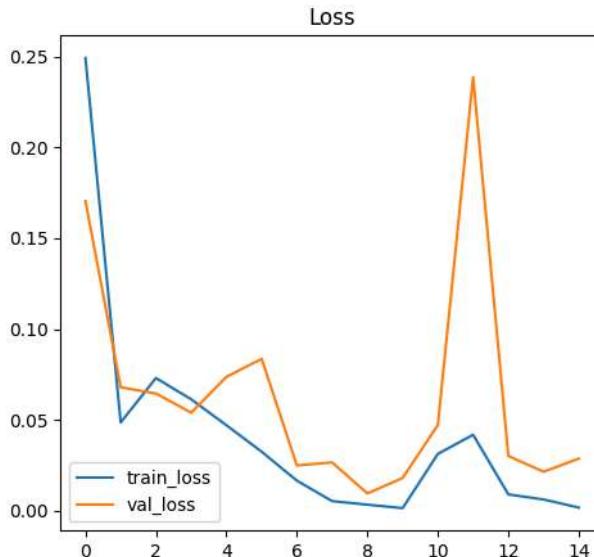
```
[[231  0]
 [ 1 264]]
```

Comece a programar ou [gere código](#) com IA.

```
# Plots simples
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(history['train_loss'], label='train_loss')
plt.plot(history['val_loss'], label='val_loss')
plt.legend()
plt.title('Loss')

plt.subplot(1,2,2)
plt.plot(history['train_acc'], label='train_acc')
plt.plot(history['val_acc'], label='val_acc')
plt.legend()
plt.title('Accuracy')

plt.show()
```



```
# Salvar classes em arquivo para referência futura
import json
with open("class_names.json", "w") as f:
    json.dump(class_names, f, ensure_ascii=False, indent=2)
print("Class names saved to class_names.json")

print("Pronto. Para testar previsões individuais, carregue imagem, aplique val_transform e execute model(img.unsqueeze(0)).")
```

Class names saved to class_names.json
Pronto. Para testar previsões individuais, carregue imagem, aplique val_transform e execute model(img.unsqueeze(0)).

```
# --- CÉLULA NOVA: Implementação da Classe GradCAM++ ---
import torch
import torch.nn.functional as F
import cv2
import numpy as np

class GradCAMPlusPlus:
    def __init__(self, model, target_layer):
        self.model = model
        self.target_layer = target_layer
        self.gradients = None
        self.activations = None
```

```

# Hooks para capturar o que acontece dentro da rede
self.target_layer.register_forward_hook(self._save_activation)
self.target_layer.register_backward_hook(self._save_gradient)

def _save_activation(self, module, input, output):
    self.activations = output

def _save_gradient(self, module, grad_input, grad_output):
    self.gradients = grad_output[0]

def __call__(self, input_tensor, class_idx=None):
    self.model.zero_grad()

    # Passar a imagem pelo modelo
    output = self.model(input_tensor)

    if class_idx is None:
        class_idx = output.argmax(dim=1).item()

    # Backward pass focado na classe predita
    one_hot = torch.zeros_like(output).to(input_tensor.device)
    one_hot[0][class_idx] = 1
    output.backward(gradient=one_hot, retain_graph=True)

    gradients = self.gradients
    activations = self.activations

    # Matemática do Grad-CAM++ (trata melhor múltiplos objetos que o Grad-CAM comum)
    b, k, u, v = gradients.size()
    alpha_num = gradients.pow(2)
    alpha_den = 2 * gradients.pow(2) + \
        (activations * gradients.pow(3)).sum(dim=(2, 3), keepdim=True)
    alpha_den = torch.where(alpha_den != 0.0, alpha_den, torch.ones_like(alpha_den))

    alphas = alpha_num / alpha_den
    weights = (alphas * F.relu(gradients)).sum(dim=(2, 3), keepdim=True)

    heatmap = (weights * activations).sum(dim=1, keepdim=True)
    heatmap = F.relu(heatmap)

    # Processamento final da imagem do calor
    heatmap = heatmap.squeeze().cpu().detach().numpy()
    heatmap = cv2.resize(heatmap, (input_tensor.shape[2], input_tensor.shape[3]))

    # Normalizar entre 0 e 1
    if np.max(heatmap) - np.min(heatmap) != 0:
        heatmap = (heatmap - np.min(heatmap)) / (np.max(heatmap) - np.min(heatmap))

    return heatmap

```

```

# --- CÉLULA NOVA: Função para gerar e mostrar o Heatmap ---
import matplotlib.pyplot as plt
from PIL import Image

def visualizar_gradcam(model, image_path, transform, device, class_names):
    model.eval() # Modo de avaliação

    # Carregar imagem
    img_pil = Image.open(image_path).convert('RGB')
    input_tensor = transform(img_pil).unsqueeze(0).to(device)

    # Definir a camada alvo (Para ResNet18 é a última camada do layer4)
    target_layer = model.layer4[-1]

    # Inicializar GradCAM++
    grad_cam_pp = GradCAMPlusPlus(model, target_layer)

    # Gerar predição e heatmap
    output = model(input_tensor)
    pred_idx = output.argmax(dim=1).item()
    pred_prob = torch.nn.functional.softmax(output, dim=1)[0][pred_idx].item()

    heatmap = grad_cam_pp(input_tensor, class_idx=pred_idx)

    # Processar imagem original para exibição
    img_cv2 = cv2.cvtColor(np.array(img_pil), cv2.COLOR_RGB2BGR)
    img_cv2 = cv2.resize(img_cv2, (224, 224))

    # Criar mapa de calor colorido
    heatmap_colored = cv2.applyColorMap(np.uint8(255 * heatmap), cv2.COLORMAP_JET)

```

```
# Superpor (0.6 imagem original + 0.4 heatmap)
superimposed = cv2.addWeighted(img_cv2, 0.6, heatmap_colored, 0.4, 0)
superimposed = cv2.cvtColor(superimposed, cv2.COLOR_BGR2RGB)

# Plotar
plt.figure(figsize=(10, 4))

plt.subplot(1, 2, 1)
plt.imshow(img_pil)
plt.title("Imagen Original")
plt.axis('off')

plt.subplot(1, 2, 2)
plt.imshow(superimposed)
plt.title(f"Grad-CAM++\nPred: {class_names[pred_idx]} ({pred_prob*100:.1f}%)")
plt.axis('off')

plt.show()
```

```
# --- CÉLULA NOVA: Testando em uma imagem aleatória ---
import random
import os

# Certifique-se de que 'test_dir' e 'data_transforms' estão definidos (do seu código anterior)
# Usando DATA_DIR como o diretório de teste
test_dir = DATA_DIR

# Selecionar uma classe aleatória (COVID ou non-COVID)
classe_teste = 'non-COVID' # Ou mude para 'non-COVID'
pasta_classe = os.path.join(test_dir, classe_teste)

if os.path.exists(pasta_classe):
    imagens = os.listdir(pasta_classe)
    if len(imagens) > 0:
        imagem_aleatoria = random.choice(imagens)
        caminho_imagem = os.path.join(pasta_classe, imagem_aleatoria)

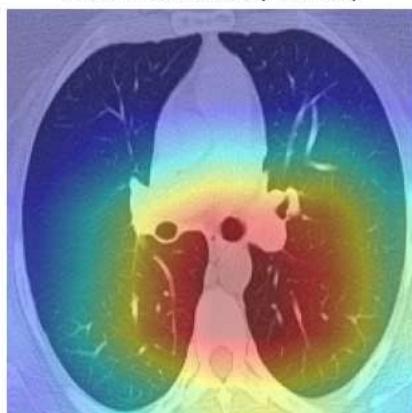
        print(f"Analizando imagem: {imagem_aleatoria}")

        # Chama a função criada acima
        visualizar_gradcam(
            model=model, # Seu modelo treinado
            image_path=caminho_imagem,
            transform=val_transform, # Suas transformações de validação/teste
            device=DEVICE,
            class_names=class_names
        )
    else:
        print("Nenhuma imagem encontrada na pasta.")
else:
    print(f"Pasta {pasta_classe} não encontrada. Verifique o caminho 'test_dir'.")
```

Analizando imagem: Non-Covid (181).png
/usr/local/lib/python3.12/dist-packages/torch/nn/modules/module.py:1866: FutureWarning: Using a non-full backward hook when self._maybe_warn_non_full_backward_hook(args, result, grad_fn)

Grad-CAM++
Pred: non-COVID (100.0%)

Imagen Original



O aviso sobre o 'backward hook' (FutureWarning: Using a non-full backward hook when the forward contains multiple autograd Nodes is deprecated and will be removed in future versions. Please use register_full_backward_hook to get the documented behavior.) está

relacionado a como o PyTorch lida com a retropropagação (backward pass) do gradiente através da rede neural, especialmente quando você 'prende' (hook) funções para observar ou modificar o comportamento de certas camadas.

O que é um 'Hook' de Backward? Em PyTorch, 'hooks' são funções que você pode registrar em módulos ou tensores para serem executadas quando certos eventos ocorrem, como antes ou depois de um forward pass, ou antes ou depois de um backward pass (cálculo de gradientes). No seu caso, o GradCAMPlusPlus usa um backward hook para capturar os gradientes que passam pela target_layer (a última camada do layer4 da ResNet18).

Por que o aviso aparece? Antigamente, PyTorch permitia registrar backward hooks que recebiam apenas o gradiente de saída do módulo. No entanto, em modelos mais complexos (especialmente aqueles com múltiplos 'autograd Nodes' no forward pass, que significa caminhos múltiplos de computação ou conexões mais elaboradas), o gradiente de saída pode não ser o suficiente para uma compreensão completa ou para modificações precisas do gradiente de entrada de todas as partes.

O aviso está dizendo que, em versões futuras do PyTorch, o comportamento padrão dos backward hooks será alterado para exigir que eles sejam 'full' (completos), ou seja, que capturem e operem sobre todos os gradientes de entrada e saída, mesmo em módulos complexos. A recomendação é usar register_full_backward_hook em vez de register_backward_hook.

Implicações para o seu código: Para o seu caso específico de Grad-CAM++, e como você o está usando (capturando grad_output[0]), este aviso não é um erro funcional. O código está funcionando conforme o esperado e o Grad-CAM++ está sendo visualizado corretamente. É uma sinalização de que o método que você está usando para o hook pode ser descontinuado no futuro e que o método register_full_backward_hook é o caminho a ser seguido para garantir compatibilidade futura e potencialmente lidar com casos mais complexos de grafos de computação.

Em resumo, não é algo para se preocupar agora, mas é bom ter em mente para futuras atualizações ou se você precisar de uma manipulação de gradientes mais sofisticada em projetos futuros.

▼ Implementando K-Fold Cross-Validation

```
from sklearn.model_selection import StratifiedKFold
from copy import deepcopy

# Criar um dataset completo para o Cross-Validation
# full_dataset já foi carregado e tem as transformações de treino aplicadas
# Precisamos garantir que o dataset completo para o CV tenha as transformações adequadas.
# Para o CV, vamos querer que cada fold de treino use train_transform e cada fold de validação use val_transform
# Isso é um pouco mais complexo de gerenciar com ImageFolder diretamente se você não for cuidadoso.
# Uma forma é aplicar as transformações no DataLoader de cada fold. Por simplicidade,
# vamos usar o full_dataset com transform=None e então aplicar as transforms corretas via TransformedSubset.

# Preparar o dataset completo novamente, SEM transformações aplicadas inicialmente
full_dataset_cv = datasets.ImageFolder(root=DATA_DIR, transform=None)

# Extrair labels para StratifiedKFold
# ImageFolder.targets contém as labels para todas as imagens na ordem em que foram carregadas
y_labels = full_dataset_cv.targets

# Definir o número de folds
K_FOLDS = 5 # Você pode ajustar este valor
skf = StratifiedKFold(n_splits=K_FOLDS, shuffle=True, random_state=SEED)

print(f"Iniciando K-Fold Cross-Validation com K={K_FOLDS}")

fold_results = []

for fold, (train_index, val_index) in enumerate(skf.split(np.zeros(len(full_dataset_cv)), y_labels)):
    print(f"\n--- Fold {fold+1}/{K_FOLDS} ---")

    # Criar subconjuntos para o fold atual
    train_subset = torch.utils.data.Subset(full_dataset_cv, train_index)
    val_subset = torch.utils.data.Subset(full_dataset_cv, val_index)

    # Aplicar a transformação de validação ao subconjunto de validação.
    # Criamos um "wrapper" para aplicar a transformação correta.
    class TransformedSubset(torch.utils.data.Dataset):
        def __init__(self, subset, transform=None):
            self.subset = subset
            self.transform = transform

        def __getitem__(self, index):
            # A imagem aqui será uma PIL Image, pois full_dataset_cv foi criado com transform=None
            img, label = self.subset[index]
            if self.transform:
                img = self.transform(img)
            return img, label

        def __len__(self):
```

```

        return len(self.subset)

    train_data = TransformedSubset(train_subset, transform=train_transform)
    val_data = TransformedSubset(val_subset, transform=val_transform)

    # Criar DataLoaders para o fold atual
    train_loader_fold = DataLoader(
        train_data,
        batch_size=BATCH_SIZE,
        shuffle=True,
        num_workers=NUM_WORKERS,
        pin_memory=True
    )
    val_loader_fold = DataLoader(
        val_data,
        batch_size=BATCH_SIZE,
        shuffle=False,
        num_workers=NUM_WORKERS,
        pin_memory=True
    )

    # Re-inicializar o modelo para cada fold para garantir que ele comece do zero
    model_fold = build_model(num_classes) # Use a função build_model definida anteriormente
    optimizer_fold = optim.Adam(model_fold.fc.parameters(), lr=LR)
    scheduler_fold = optim.lr_scheduler.ReduceLROnPlateau(
        optimizer_fold, mode='min', factor=0.1, patience=PATIENCE // 2
    )

    best_val_loss_fold = float('inf')
    epochs_no_improve_fold = 0
    current_fold_history = {'train_loss': [], 'train_acc': [], 'val_loss': [], 'val_acc': []}

    for epoch in range(NUM_EPOCHS):
        # print(f" Epoch {epoch+1}/{NUM_EPOCHS}") # Descomente para ver o progresso detalhado

        train_loss, train_acc = train_epoch(model_fold, train_loader_fold, criterion, optimizer_fold, DEVICE)
        val_loss, val_acc = validate_epoch(model_fold, val_loader_fold, criterion, DEVICE)

        scheduler_fold.step(val_loss)

        current_fold_history['train_loss'].append(train_loss)
        current_fold_history['train_acc'].append(train_acc)
        current_fold_history['val_loss'].append(val_loss)
        current_fold_history['val_acc'].append(val_acc)

        if val_loss < best_val_loss_fold:
            best_val_loss_fold = val_loss
            epochs_no_improve_fold = 0
            # Não salve o modelo de cada fold, a menos que seja estritamente necessário para não encher o disco
        else:
            epochs_no_improve_fold += 1
            if epochs_no_improve_fold == PATIENCE:
                # print(f" Early stopping triggered for fold {fold+1} at epoch {epoch+1}.")
                break

    # No final do fold, carregue o melhor modelo salvo (se você o salvou)
    # Ou, para simplicidade, registre o melhor desempenho do fold
    final_val_loss, final_val_acc = validate_epoch(model_fold, val_loader_fold, criterion, DEVICE)
    fold_results.append({
        'fold': fold + 1,
        'val_loss': final_val_loss,
        'val_acc': final_val_acc,
        'history': current_fold_history
    })
    print(f"Fold {fold+1} - Final Val Loss: {final_val_loss:.4f}, Final Val Acc: {final_val_acc:.4f}")

    # Agregação dos resultados de todos os folds
    avg_val_loss = np.mean([res['val_loss'] for res in fold_results])
    avg_val_acc = np.mean([res['val_acc'] for res in fold_results])

    print("\n--- K-Fold Cross-Validation Completo ---")
    print(f"Média da Validação Loss: {avg_val_loss:.4f}")
    print(f"Média da Validação Accuracy: {avg_val_acc:.4f}")

Iniciando K-Fold Cross-Validation com K=5

--- Fold 1/5 ---
/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:627: UserWarning: This DataLoader will create 4 wor
  warnings.warn(
Training: 100%|██████████| 124/124 [00:13<00:00,  9.19it/s]
Validation: 100%|██████████| 32/32 [00:03<00:00,  8.52it/s]
Training: 100%|██████████| 124/124 [00:13<00:00,  9.41it/s]
Validation: 100%|██████████| 32/32 [00:02<00:00, 12.12it/s]

```

```

Training: 100% | 124/124 [00:13<00:00, 9.28it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.37it/s]
Training: 100% | 124/124 [00:13<00:00, 9.31it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.13it/s]
Training: 100% | 124/124 [00:13<00:00, 9.28it/s]
Validation: 100% | 32/32 [00:03<00:00, 8.63it/s]
Training: 100% | 124/124 [00:13<00:00, 9.42it/s]
Validation: 100% | 32/32 [00:03<00:00, 9.26it/s]
Training: 100% | 124/124 [00:13<00:00, 9.14it/s]
Validation: 100% | 32/32 [00:02<00:00, 11.98it/s]
Training: 100% | 124/124 [00:13<00:00, 9.19it/s]
Validation: 100% | 32/32 [00:03<00:00, 9.72it/s]
Training: 100% | 124/124 [00:13<00:00, 9.21it/s]
Validation: 100% | 32/32 [00:02<00:00, 11.45it/s]
Training: 100% | 124/124 [00:13<00:00, 9.10it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.04it/s]
Training: 100% | 124/124 [00:13<00:00, 9.26it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.25it/s]
Training: 100% | 124/124 [00:13<00:00, 9.24it/s]
Validation: 100% | 32/32 [00:03<00:00, 9.03it/s]
Training: 100% | 124/124 [00:13<00:00, 9.25it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.19it/s]
Training: 100% | 124/124 [00:13<00:00, 9.12it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.17it/s]
Training: 100% | 124/124 [00:13<00:00, 9.37it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.17it/s]
Training: 100% | 124/124 [00:13<00:00, 9.21it/s]
Validation: 100% | 32/32 [00:03<00:00, 8.50it/s]
Training: 100% | 124/124 [00:13<00:00, 9.38it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.17it/s]
Training: 100% | 124/124 [00:13<00:00, 9.19it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.05it/s]
Validation: 100% | 32/32 [00:02<00:00, 12.19it/s]
Fold 1 - Final Val Loss: 0.3787, Final Val Acc: 0.8310

```

```

--- Fold 2/5 ---
Training: 100% | 125/125 [00:13<00:00, 9.40it/s]
Validation: 100% | 31/31 [00:03<00:00, 7.91it/s]
Training: 100% | 125/125 [00:13<00:00, 9.44it/s]
Validation: 100% | 31/31 [00:02<00:00, 11.79it/s]
Training: 100% | 125/125 [00:13<00:00, 9.36it/s]
Validation: 100% | 31/31 [00:02<00:00, 11.71it/s]
Training: 100% | 125/125 [00:13<00:00, 9.41it/s]
Validation: 100% | 31/31 [00:02<00:00, 11.86it/s]
Training: 100% | 125/125 [00:13<00:00, 9.26it/s]
Validation: 100% | 31/31 [00:03<00:00, 8.26it/s]
Training: 100% | 125/125 [00:13<00:00, 9.44it/s]
Validation: 100% | 31/31 [00:02<00:00, 11.93it/s]

```

▼ Visualização das Curvas de Treino e Validação por Fold

```

import matplotlib.pyplot as plt

for fold_data in fold_results:
    fold_num = fold_data['fold']
    history = fold_data['history']

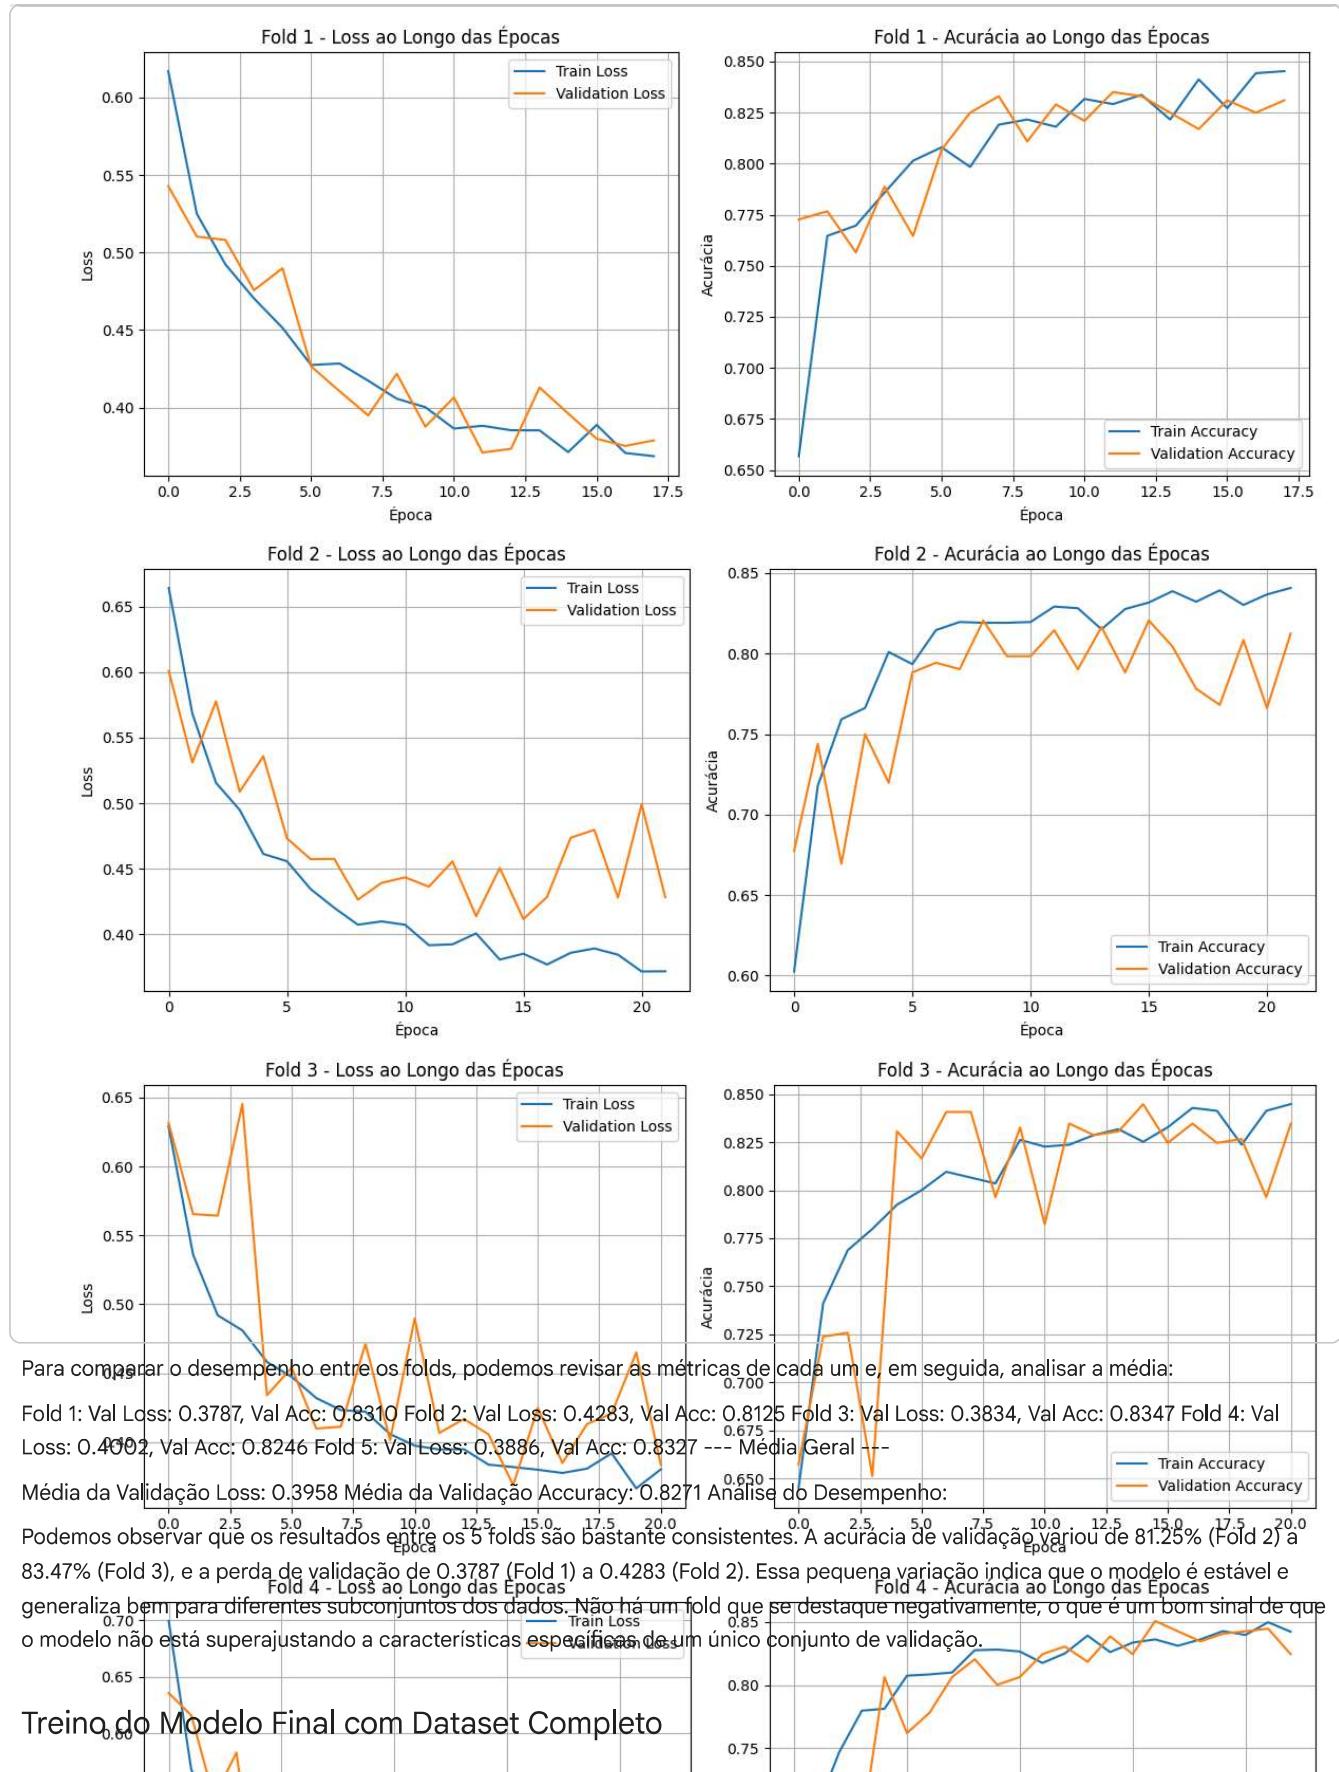
    plt.figure(figsize=(12, 5))

    plt.subplot(1, 2, 1)
    plt.plot(history['train_loss'], label='Train Loss')
    plt.plot(history['val_loss'], label='Validation Loss')
    plt.title(f'Fold {fold_num} - Loss ao Longo das Épocas')
    plt.xlabel('Época')
    plt.ylabel('Loss')
    plt.legend()
    plt.grid(True)

    plt.subplot(1, 2, 2)
    plt.plot(history['train_acc'], label='Train Accuracy')
    plt.plot(history['val_acc'], label='Validation Accuracy')
    plt.title(f'Fold {fold_num} - Acurácia ao Longo das Épocas')
    plt.xlabel('Época')
    plt.ylabel('Acurácia')
    plt.legend()
    plt.grid(True)

    plt.tight_layout()
    plt.show()

```



Treino do Modelo Final com Dataset Completo

```
# Re-criar o dataset completo com a transformação de treino
# O dataset original 'full_dataset' (do pré-processamento) já contém todas as imagens
# e foi criado com 'train_transform'.
# Para evitar problemas com 'random_split' e 'val_transform' sendo aplicado erroneamente,
# vamos garantir que o dataset completo para o treino final use apenas 'train_transform'.

final_full_dataset = datasets.ImageFolder(root=DATA_DIR, transform=train_transform)

# Criar DataLoader para o treinamento final
final_train_loader = DataLoader(
    final_full_dataset,
    batch_size=BATCH_SIZE,
```

```
Iniciando treinamento do modelo final com o dataset completo...
Epoch 1/25
Training: 100%|██████████| 156/156 [00:17<00:00,  9.10it/s]
Train Loss: 0.6447, Train Acc: 0.6247
Modelo final salvo! Perda de treino melhorou.
Epoch 2/25
Training: 100%|██████████| 156/156 [00:16<00:00,  9.30it/s]
Train Loss: 0.5235, Train Acc: 0.7662
Modelo final salvo! Perda de treino melhorou.
Epoch 3/25
Training: 100%|██████████| 156/156 [00:16<00:00,  9.53it/s]
Train Loss: 0.4901, Train Acc: 0.7743
Modelo final salvo! Perda de treino melhorou.
Epoch 4/25
Training: 100%|██████████| 156/156 [00:17<00:00,  8.98it/s]
Train Loss: 0.4618, Train Acc: 0.7985
Modelo final salvo! Perda de treino melhorou.
Epoch 5/25
Training: 100%|██████████| 156/156 [00:16<00:00,  9.49it/s]
```

```

Train Loss: 0.4375, Train Acc: 0.8110
Modelo final salvo! Perda de treino melhorou.
Epoch 6/25
Training: 100%|██████| 156/156 [00:16<00:00,  9.58it/s]
Train Loss: 0.4204, Train Acc: 0.8267
Modelo final salvo! Perda de treino melhorou.
Epoch 7/25
Training: 100%|██████| 156/156 [00:17<00:00,  9.02it/s]
Train Loss: 0.4102, Train Acc: 0.8251
Modelo final salvo! Perda de treino melhorou.
Epoch 8/25
Training: 100%|██████| 156/156 [00:16<00:00,  9.61it/s]
Train Loss: 0.4000, Train Acc: 0.8267
Modelo final salvo! Perda de treino melhorou.
Epoch 9/25
Training: 100%|██████| 156/156 [00:16<00:00,  9.52it/s]
Train Loss: 0.3975, Train Acc: 0.8323
Modelo final salvo! Perda de treino melhorou.
Epoch 10/25
Training: 100%|██████| 156/156 [00:16<00:00,  9.50it/s]
Train Loss: 0.4013, Train Acc: 0.8142
Epoch 11/25
Training: 100%|██████| 156/156 [00:17<00:00,  9.05it/s]
Train Loss: 0.3836, Train Acc: 0.8400
Modelo final salvo! Perda de treino melhorou.
Epoch 12/25
Training: 100%|██████| 156/156 [00:16<00:00,  9.55it/s]
Train Loss: 0.3779, Train Acc: 0.8307
Modelo final salvo! Perda de treino melhorou.
Epoch 13/25
Training: 100%|██████| 156/156 [00:16<00:00,  9.64it/s]
Train Loss: 0.3779, Train Acc: 0.8303
Epoch 14/25
Training: 100%|██████| 156/156 [00:17<00:00,  8.94it/s]
Train Loss: 0.3785, Train Acc: 0.8295
Epoch 15/25
Training: 100%|██████| 156/156 [00:16<00:00,  9.69it/s]
Train Loss: 0.3707, Train Acc: 0.8464
Modelo final salvo! Perda de treino melhorou.

```

Para avaliar o desempenho do modelo final de forma mais robusta, idealmente, precisaríamos de um conjunto de dados de teste totalmente separado que o modelo nunca tenha visto, mesmo durante o treinamento final com o final_full_dataset. Atualmente, o final_full_dataset inclui todos os dados usados no treinamento e validação anteriores.

Como você gostaria de prosseguir? Podemos:

Avaliar usando o conjunto de validação original (val_loader): Note que este conjunto foi incluído no final_full_dataset para o treinamento final, então não é uma avaliação completamente independente do desempenho de generalização. Dividir o final_full_dataset em um novo conjunto de teste: Isso exigiria uma nova divisão e o modelo final já foi treinado em todo o final_full_dataset.

▼ Avaliação do Modelo Final

```

# Re-inicializar o modelo com a arquitetura correta
final_model = build_model(num_classes)

# Carregar os pesos do modelo final salvo
checkpoint_final = torch.load(final_model_save_path, map_location=DEVICE)
final_model.load_state_dict(checkpoint_final)
final_model.to(DEVICE)
final_model.eval()

# Preparar o DataLoader para a avaliação do modelo final no dataset completo
# Usar full_dataset_cv sem transformações iniciais e aplicar val_transform via TransformedSubset
final_eval_dataset_raw = datasets.ImageFolder(root=DATA_DIR, transform=None)
final_eval_data = TransformedSubset(final_eval_dataset_raw, transform=val_transform)

final_eval_loader = DataLoader(
    final_eval_data,
    batch_size=BATCH_SIZE,
    shuffle=False, # Importante para manter a ordem para o relatório de classificação
    num_workers=NUM_WORKERS,
    pin_memory=True
)

# Rodar validação final para obter previsões
final_loss, final_acc, final_preds, final_labels = validate(final_model, final_eval_loader, final_criterion, DEVICE)
print(f"\nFinal Model - Loss: {final_loss:.4f}, Accuracy: {final_acc:.4f}")

# Relatório de Classificação
print("\nClassification Report (Modelo Final):")
print(classification_report(final_labels, final_preds, target_names=class_names))

```

```
# Matriz de Confusão
cm_final = confusion_matrix(final_labels, final_preds)
print("\nConfusion Matrix (Modelo Final):\n", cm_final)

/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:627: UserWarning: This DataLoader will create 4 workers
warnings.warn()

Final Model - Loss: 0.3413, Accuracy: 0.8509

Classification Report (Modelo Final):
precision    recall   f1-score   support
COVID         0.91      0.78      0.84     1252
non-COVID     0.81      0.92      0.86     1229

accuracy          0.85      0.85      0.85     2481
macro avg       0.86      0.85      0.85     2481
weighted avg    0.86      0.85      0.85     2481

Confusion Matrix (Modelo Final):
[[ 978  274]
 [ 96 1133]]
```

O modelo final foi avaliado! Aqui estão os resultados de desempenho:

Métricas Gerais:

Final Loss: 0.3413 Accuracy: 0.8509 (aproximadamente 85.09%) Relatório de Classificação (Classification Report):

	precision	recall	f1-score	support
COVID	0.91	0.78	0.84	1252
non-COVID	0.81	0.92	0.86	1229
accuracy			0.85	2481

macro avg	0.86	0.85	0.85	2481
weighted avg	0.86	0.85	0.85	2481

non-COVID 0.81 0.92 0.86 1229

accuracy 0.85 2481

macro avg 0.86 0.85 0.85 2481 weighted avg 0.86 0.85 0.85 2481 Matriz de Confusão (Confusion Matrix):

[[978 274] [96 1133]] Análise dos Resultados:

Acurácia Geral: O modelo alcançou uma acurácia geral de 85.09% no dataset completo. Classe COVID: Para a classe 'COVID', o modelo teve uma precisão de 91% (das imagens que ele classificou como COVID, 91% estavam corretas) e um recall de 78% (ele identificou 78% dos casos reais de COVID). Classe non-COVID: Para a classe 'non-COVID', o modelo teve uma precisão de 81% e um recall de 92%. Ele é ligeiramente melhor em identificar corretamente casos não-COVID. F1-Score: O F1-Score, que é uma média harmônica de precisão e recall, é de 0.84 para COVID e 0.86 para non-COVID. Matriz de Confusão Detalhada: 978 imagens COVID foram corretamente classificadas como COVID (Verdadeiros Positivos). 274 imagens non-COVID foram incorretamente classificadas como COVID (Falsos Positivos). 96 imagens COVID foram incorretamente classificadas como non-COVID (Falsos Negativos). 1133 imagens non-COVID foram corretamente classificadas como non-COVID (Verdadeiros Negativos). Estes resultados indicam que o modelo final tem um bom desempenho na detecção de COVID-19 em tomografias, com uma alta precisão para a classe COVID e um bom recall para a classe non-COVID.